

# Operating System Concepts

---

Tei-Wei Kuo

[ktw@csie.ntu.edu.tw](mailto:ktw@csie.ntu.edu.tw)

Dept. of Computer Science and  
Information Engineering  
National Taiwan University

# Contents

1. Introduction
2. System Structures
3. Process Concept
4. Multithreaded Programming
5. Process Scheduling
6. Synchronization
7. Deadlocks
8. Memory-Management Strategies
9. Virtual-Memory Management
10. File System
11. Mass-Storage Structures
12. I/O Systems
13. Protection, Security, Distributed Systems 2

# Chapter 1. Introduction

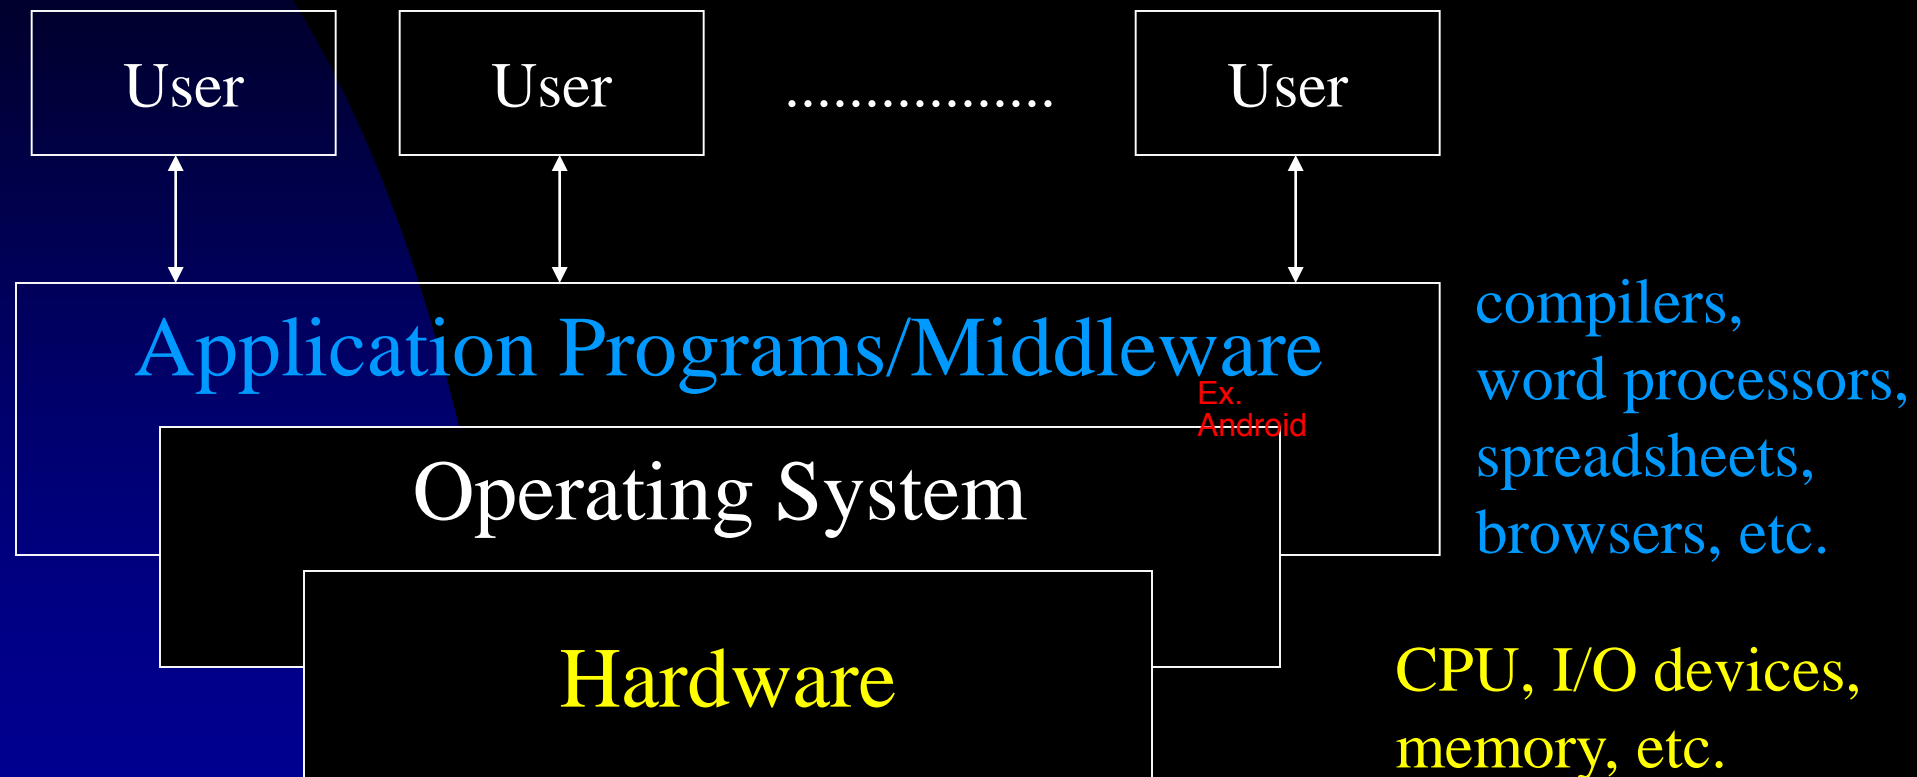
# Introduction

- What is an Operating System?
  - A basis for application programs
  - An intermediary between users and hardware
- Amazing variety
  - Mainframe, personal computer (PC), handheld computer, embedded computer without any user view

大型主機

Convenient vs Efficient

# Computer System Components



- OS – a government/environment provider

# User View

- The user view of the computer varies by the interface being used!
- Examples:
  - Personal Computer → Ease of use
  - Mainframe or minicomputer → maximization of resource utilization
    - Efficiency and fair share
  - Workstations → compromise between individual usability & resource utilization
  - Handheld computer → individual usability
  - Embedded computer without user view → run without user intervention

# System View

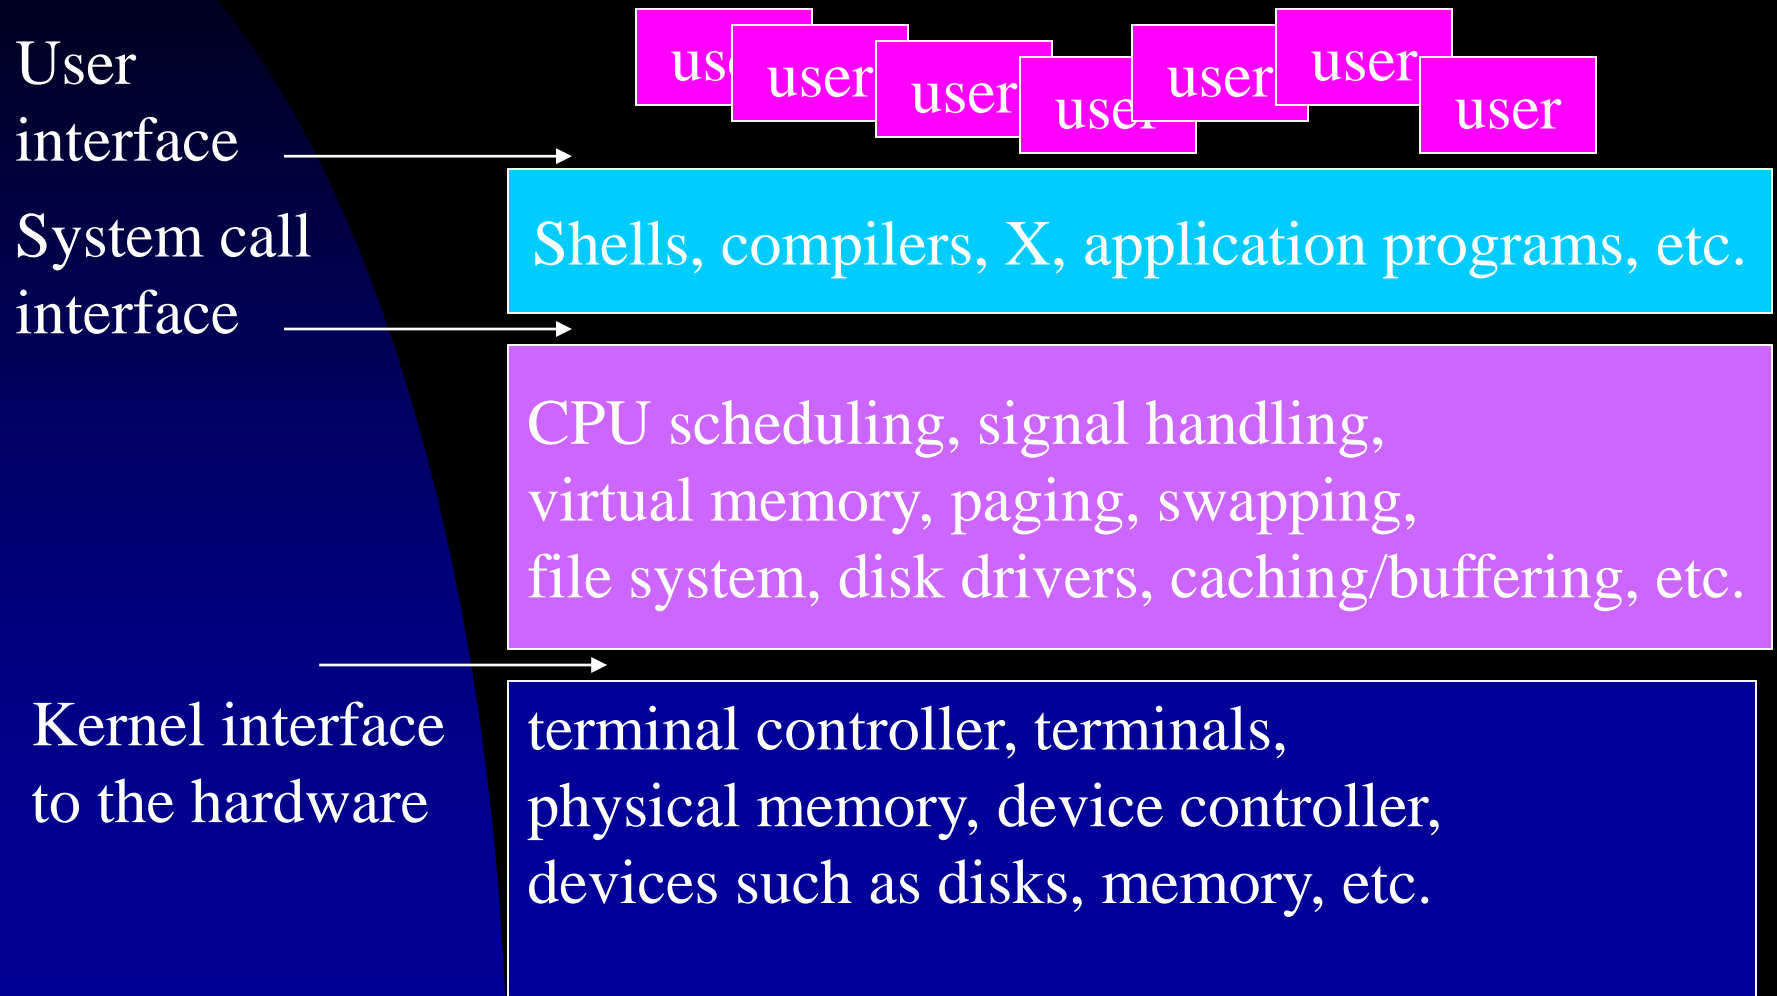
- A Resource Allocator
  - CPU time, Memory Space, File Storage, I/O Devices, Shared Code, Data Structures, and more
- A Control Program
  - Control execution of user programs
  - Prevent errors and misuse
- OS definitions – US Dept.of Justice against Microsoft in 1998
  - The stuff shipped by vendors as an OS
  - Run at all time

# System Goals

- Two Conflicting Goals:
  - Convenient for the user!
  - Efficient operation of the computer system!
- We should
  - recognize the influences of operating systems and computer architecture on each other
  - and learn why and how OS's are by tracing their evolution and predicting what they will become!



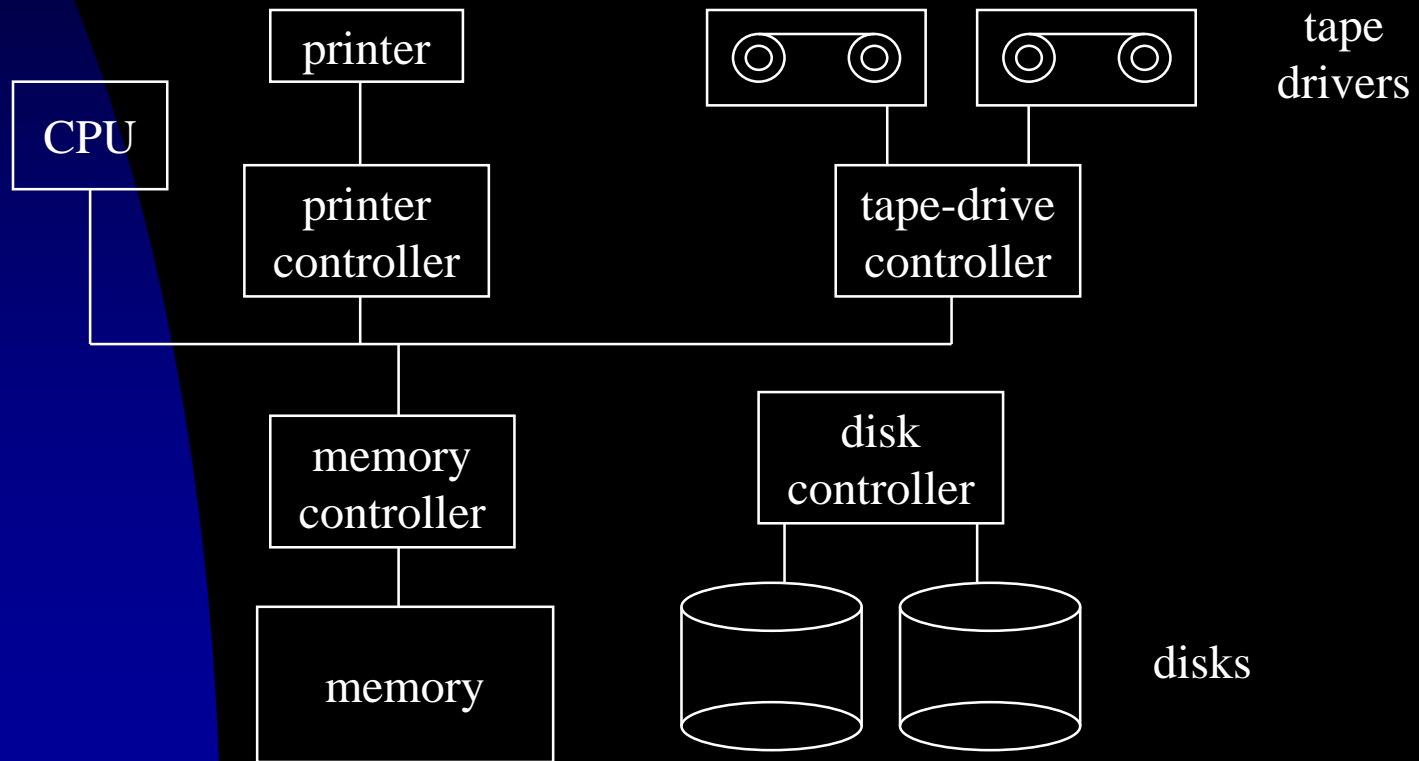
# UNIX Architecture



UNIX

# Computer-System Organization

- Objective: General knowledge of the structure of a computer system.



- Device controllers: synchronize and manage access to devices.

# Booting

電腦開啟時，需要跑一個initial program (bootstrap program) 儲存在ROM或EEPROM，(俗稱韌體firmware)

- Bootstrap program:
  - Initialize all aspects of the system, e.g., CPU registers, device controllers, memory, etc.
  - Load and run the OS
- Operating system: run *init* to initialize system processes, e.g., various daemons, login processes, after the kernel has been bootstrapped. (/etc/rc\* & init or /sbin/rc\* & init)

# Interrupt

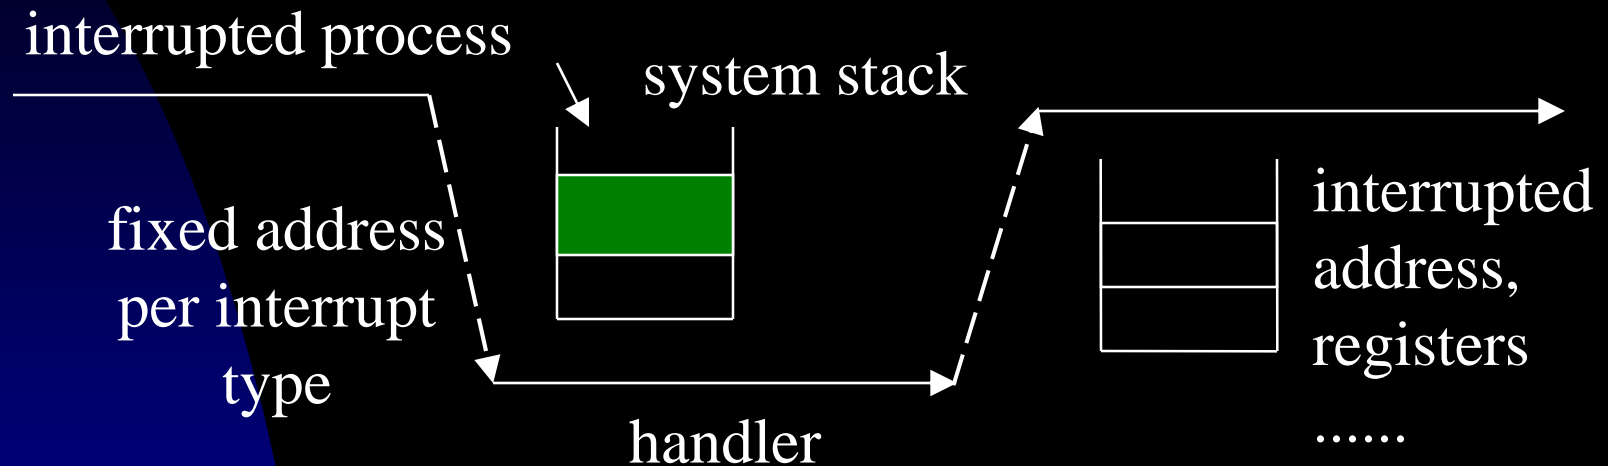
- Hardware interrupt, e.g. services requests of I/O devices
- Software interrupt, e.g. signals, invalid memory access, division by zero, system calls, etc – (trap)




- Procedures: generic handler or interrupt vector (MS-DOS, UNIX)



# Interrupt Handling Procedure

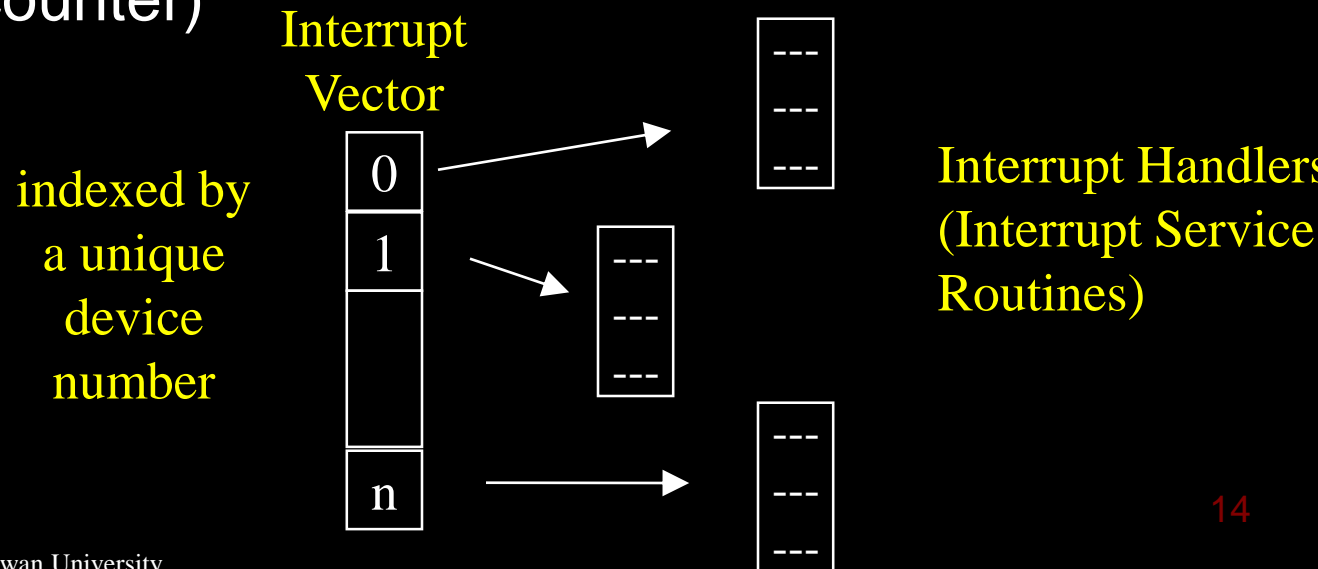


- Saving of the address of the interrupted instruction: fixed locations or stacks
- Interrupt disabling or enabling issues: lost interrupt?! 

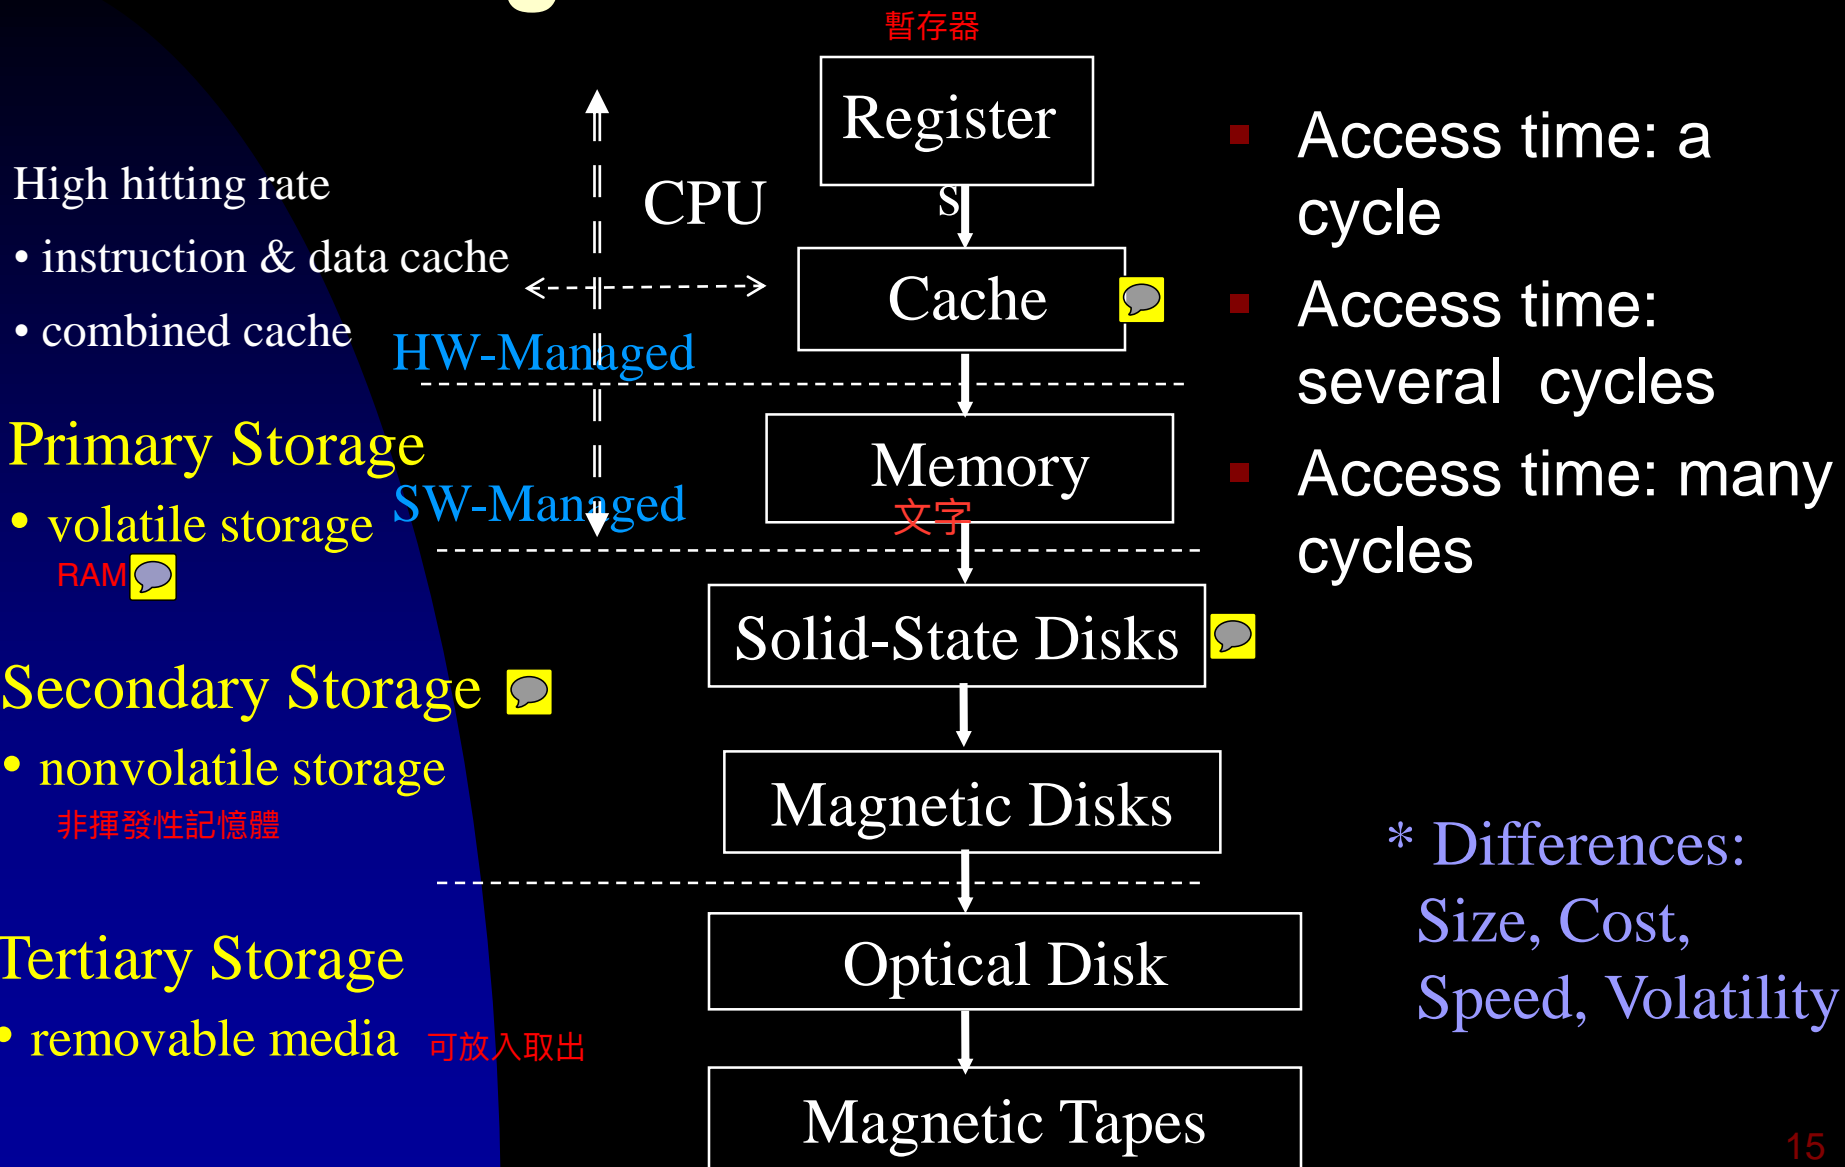
prioritized interrupts → masking

# Interrupt Handling Procedure

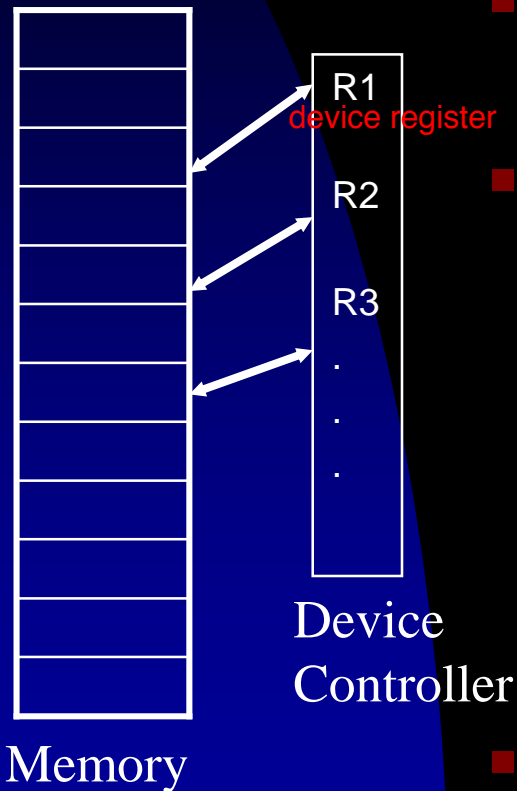
- Interrupt Handling
  - ➔ Save interrupt information
  - ➔ OS determine the interrupt type (by polling)
  - ➔ Call the corresponding handlers
  - ➔ Return to the interrupted job by the restoring important information (e.g., saved return addr. ➔ program counter)



# Storage Structure




# Memory



- Processor can have direct access!
- Intermediate storage for data in the registers of device controllers
- Memory-Mapped I/O (PC & Mac)
  - (1) Frequently used devices
  - (2) Devices must be fast, such as video controller, or special I/O instructions is used to move data between memory & device controller registers
- Programmed I/O – polling
  - or interrupt-driven handling



# Magnetic disks



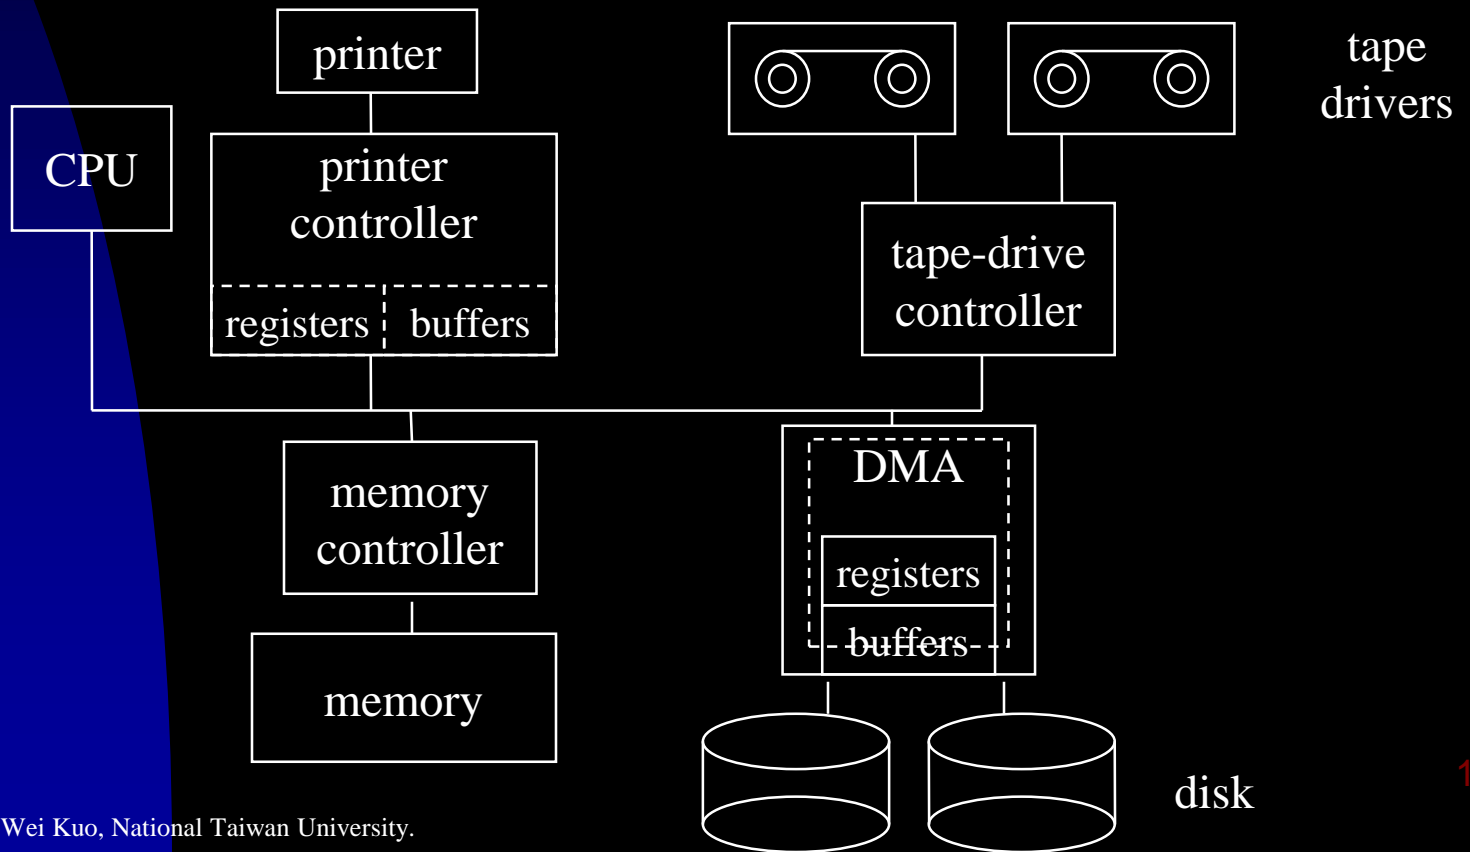
- Transfer Rate
- Random-Access Time
  - Seek time in x ms
  - Rotational latency in y ms
    - 60~200 times/sec

# Magnetic Disks

- Disks
  - Fixed-head disks:
    - More r/w heads v.s. fast track switching
  - Moving-head disks (hard disk)
  - Primary concerns:
    - Cost, Size, Speed
  - Computer → host controller → disk controller → disk drives (cache ↔ disks)
- Floppy disk
  - slow rotation, low capacity, low density, but less expensive
- Tapes: backup or data transfer bet machines

# I/O Structure

- Device controllers are responsible of moving data between the peripheral devices and their local buffer storages.



# I/O Structure

- I/O operation

- a. CPU sets up specific controller registers within the controller.

- b. Read: devices → controller buffers → **memory**

- Write: **memory** → controller buffers → devices



- c. Notify the completion of the operation by triggering an interrupt

# DMA

直接記憶體存取

一次做完I/O，再中斷，而不是每一次都中斷

- Goal: Release CPU from handling excessive interrupts!

- E.g. 9600-baud terminal



2-microsecond service / 1000 microseconds

High-speed device:

50%時間處理中  
斷

2-microsecond service / 4 microseconds

- Procedure

- Execute the device driver to set up the registers of the DMA controller.
- DMA moves blocks of data between the memory and its own buffers.
- Transfer from its buffers to its devices.
- Interrupt the CPU when the job is done.



# Single-Processor Systems

- Characteristics: One Main CPU
  - Special-Purpose Processors, e.g., Disk-Controller Microprocessors.
- Examples:
  - Personal Computers (Since 1970's), Mainframes.
- Operating Systems
  - Batching → Multiprogramming → Time-Sharing

# Multiprocessor/Parallel Systems


- Tightly coupled: have more than one processor in close communication sharing computer bus, clock, and sometimes memory and peripheral devices 外部裝置
- Loosely coupled: otherwise
- Advantages
  - Speedup – Throughput
  - Lower cost – Economy of Scale
  - More reliable – Graceful Degradation → Fail Soft (detection, diagnosis, correction)
    - A Tandem or HP-NonStop fault-tolerance solution

# Multiprocessor/Parallel Systems

- Symmetric multiprocessing model: each processor runs an identical copy of the OS
- Asymmetric multiprocessing model: a master-slave relationship
  - ~ Dynamically allocate or pre-allocate tasks
  - ~ Commonly seen in extremely large systems
  - ~ Hardware and software make a difference?
- Trend: the dropping of microporcessor cost  
→ OS functions are offloaded to slave processors (back-ends)



# Multiprocessor/Parallel Systems

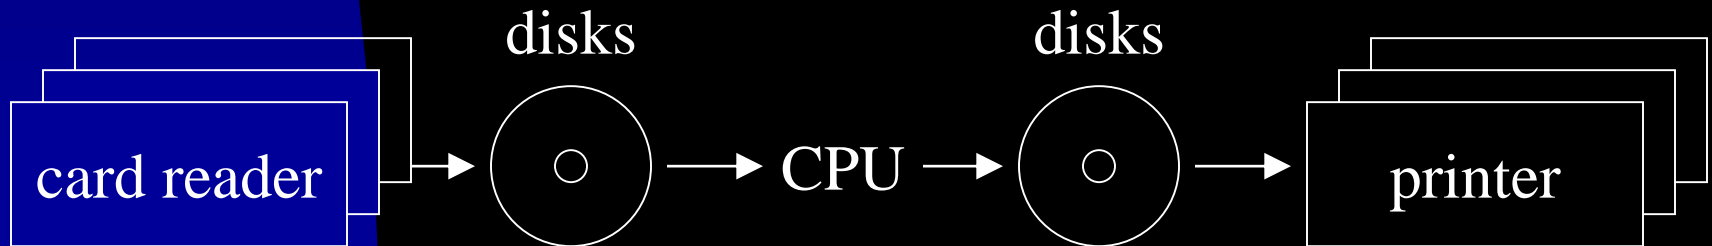
- The Recent Trend: 
  - Hyperthreading Processors
  - Multiple Cores over a Single Chip
    - N Standard Processors!
- Loosely-Coupled Systems
  - Processors do not share memory or a clock
  - Blade Servers
    - Each blade-processor board boots independently and runs its own OS.

# Clustered Systems

- Definition: Clustered computers which share storage and are closely linked via LAN networking.
- Advantages: high availability, performance improvement, etc.
- Types
  - Asymmetric/symmetric clustering
  - Parallel clustering – multiple hosts that access the same data on the shared storage.
- Distributed Lock Manager (DLM)
  - Oracle

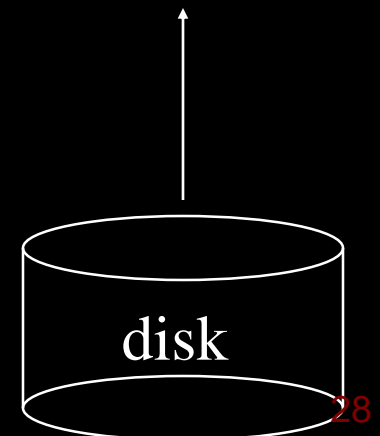
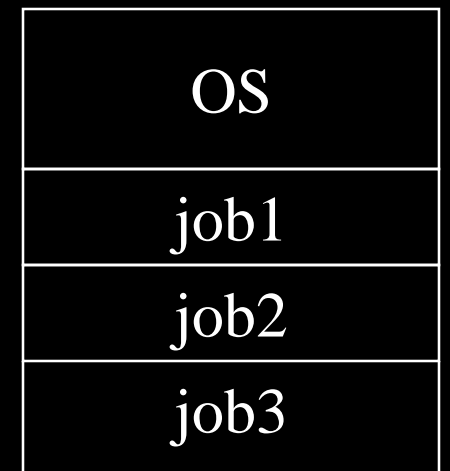
# Operating-System Structure

- Simple batch systems
  - Resident monitor – Automatically transfer control from one job to the next 工作做完後搬下一個進去
- Spooling (Simultaneous Peripheral Operation On-Line)
  - Replace sequential-access devices with random-access device



# Operating-System Structure

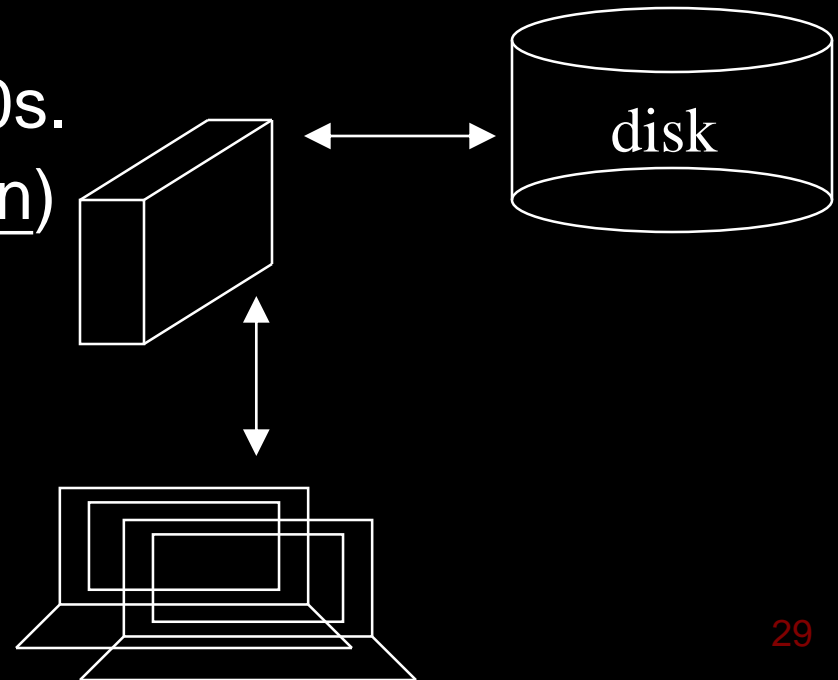
- Multiprogramming increases CPU utilization by organizing jobs so that the CPU always has one to execute – Early 1960
  - Job scheduling and CPU scheduling
  - Goal : efficient use of scarce resources



# Operating-System Structure

- Time sharing (or multitasking) is a logical extension of multiprogramming!
  - Started in 1960s and become common in 1970s.
  - An interactive (or hand-on) computer system
  - Multics, IBM OS/360
- Virtual Memory
  - Physical Address

on-line file system  
virtual memory  
sophisticated CPU scheduling  
job synchronization  
protection & security  
.....  
and so on



# Operating-System Operations

- An Interrupt-Driven Architecture for Modern OS's
  - Events are almost always signaled by the occurrence of an interrupt or a trap (or an exception).
- Protection of User Programs and OS
  - Multiprogramming
  - Sharing of Hardware and Software

# Hardware Protection

- Goal:
  - Prevent errors and misuse!
    - E.g., input errors of a program in a simple batch operating system
    - E.g., the modifications of data and code segments of another process or OS
- Dual-Mode Operations – a mode bit
  - User-mode executions except those after a trap or an interrupt occurs.
  - Monitor-mode (system mode, privileged mode, supervisor mode)
    - Privileged instruction: Machine instructions that may cause harm

# Hardware Protection

- More Modes:
  - One for the Virtual Machine Manager – It provides an interface that is identical to the underlying bare hardware.
  - More for different kernel components

virtual  
user  
mode  
virtual  
monitor  
mode  
monitor  
mode

User  
mode

processes	processes	processes
kernel 1	kernel 2	kernel 3
virtual machine software		
hardware		




# Hardware Protection

- System Calls – trap to OS for executing privileged instructions.
- Resources to protect
  - I/O devices, Memory, CPU
- I/O Protection (I/O devices are scare resources!)
  - I/O instructions are privileged.
    - User programs must issue I/O through OS
    - User programs can never gain control over the computer in the system mode.

# Hardware Protection

- Memory Protection
  - Goal: Prevent a user program from modifying the code or data structures of either the OS or other users!
  - Instructions to modify the memory space for a process are privileged.



Base register

Limit register

⇔ Check for every  
memory address by  
hardware

# Hardware Protection

- CPU Protection
  - Goal
    - Prevent user programs from sucking up CPU power!
  - Use a timer to implement time-sharing or to compute the current time.
    - Instructions that modify timers are privileged.
  - Computer control is turned over to OS for every time-slice of time!
    - Terms: time-sharing, context switch  
某個程式執行的狀態

# System Components – Process Management

- Process Management
  - Process: An Active Entity
    - Physical and Logical Resources
      - Memory, I/O buffers, data, etc.
    - Data Structures Representing Current Activities:
      - Program Counter    記錄執行到哪一行
      - Stack
      - Data Section
      - CPU Registers
      - ....
      - And More

Passive entity

Program  
(code)

+

# System Components – Process Management

- Services
  - Process creation and deletion
  - Process suspension and resumption
  - Process synchronization
  - Process communication
  - Deadlock handling

# System Components – Memory Management

- Memory: a large array of words or bytes, where each has its own address
- OS must keep several programs in memory to improve CPU utilization and user response time
- Management algorithms depend on the hardware support
- Services
  - Memory usage and availability
  - Decision of memory assignment
  - Memory allocation and deallocation

# System Components – File-System Management

- Goal:
  - A uniform logical view of information storage
  - Each medium controlled by a device
    - Magnetic tapes, magnetic disks, optical disks, etc.
- OS provides a logical storage unit: File
  - Formats:
    - Free form or being formatted rigidly.
  - General Views:
    - A sequence of bits, bytes, lines, records

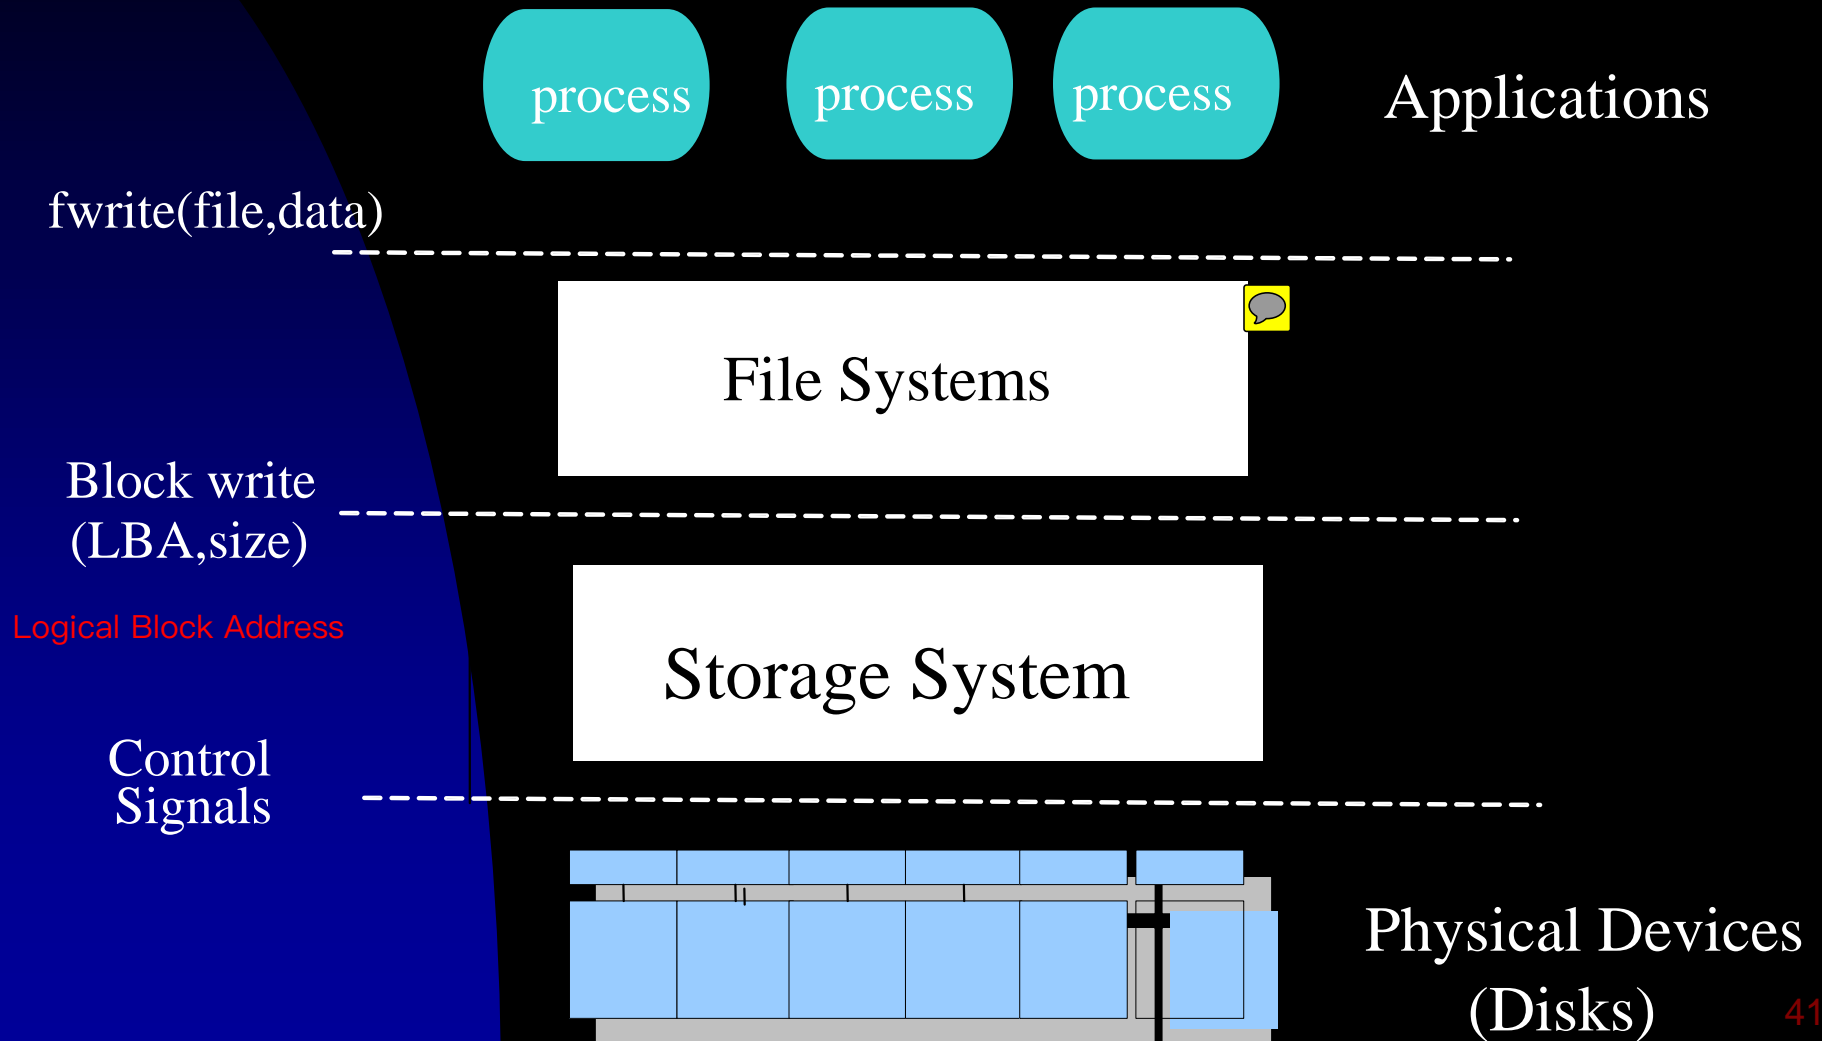
# System Components – File-System Management

- Services
  - File creation and deletion
  - Directory creation and deletion
  - Primitives for file and directory manipulation
  - Mapping of files onto secondary storage
  - File Backup

\* Privileges for file access control



# System Components – File-System Management



# System Components – Mass-Storage Management

- Goal:
  - On-line storage medium for programs & data
    - Backup of main memory
- Services for Disk Management
  - Free-space management
  - Storage allocation, e.g., continuous allocation
  - Disk scheduling, e.g., FCFS

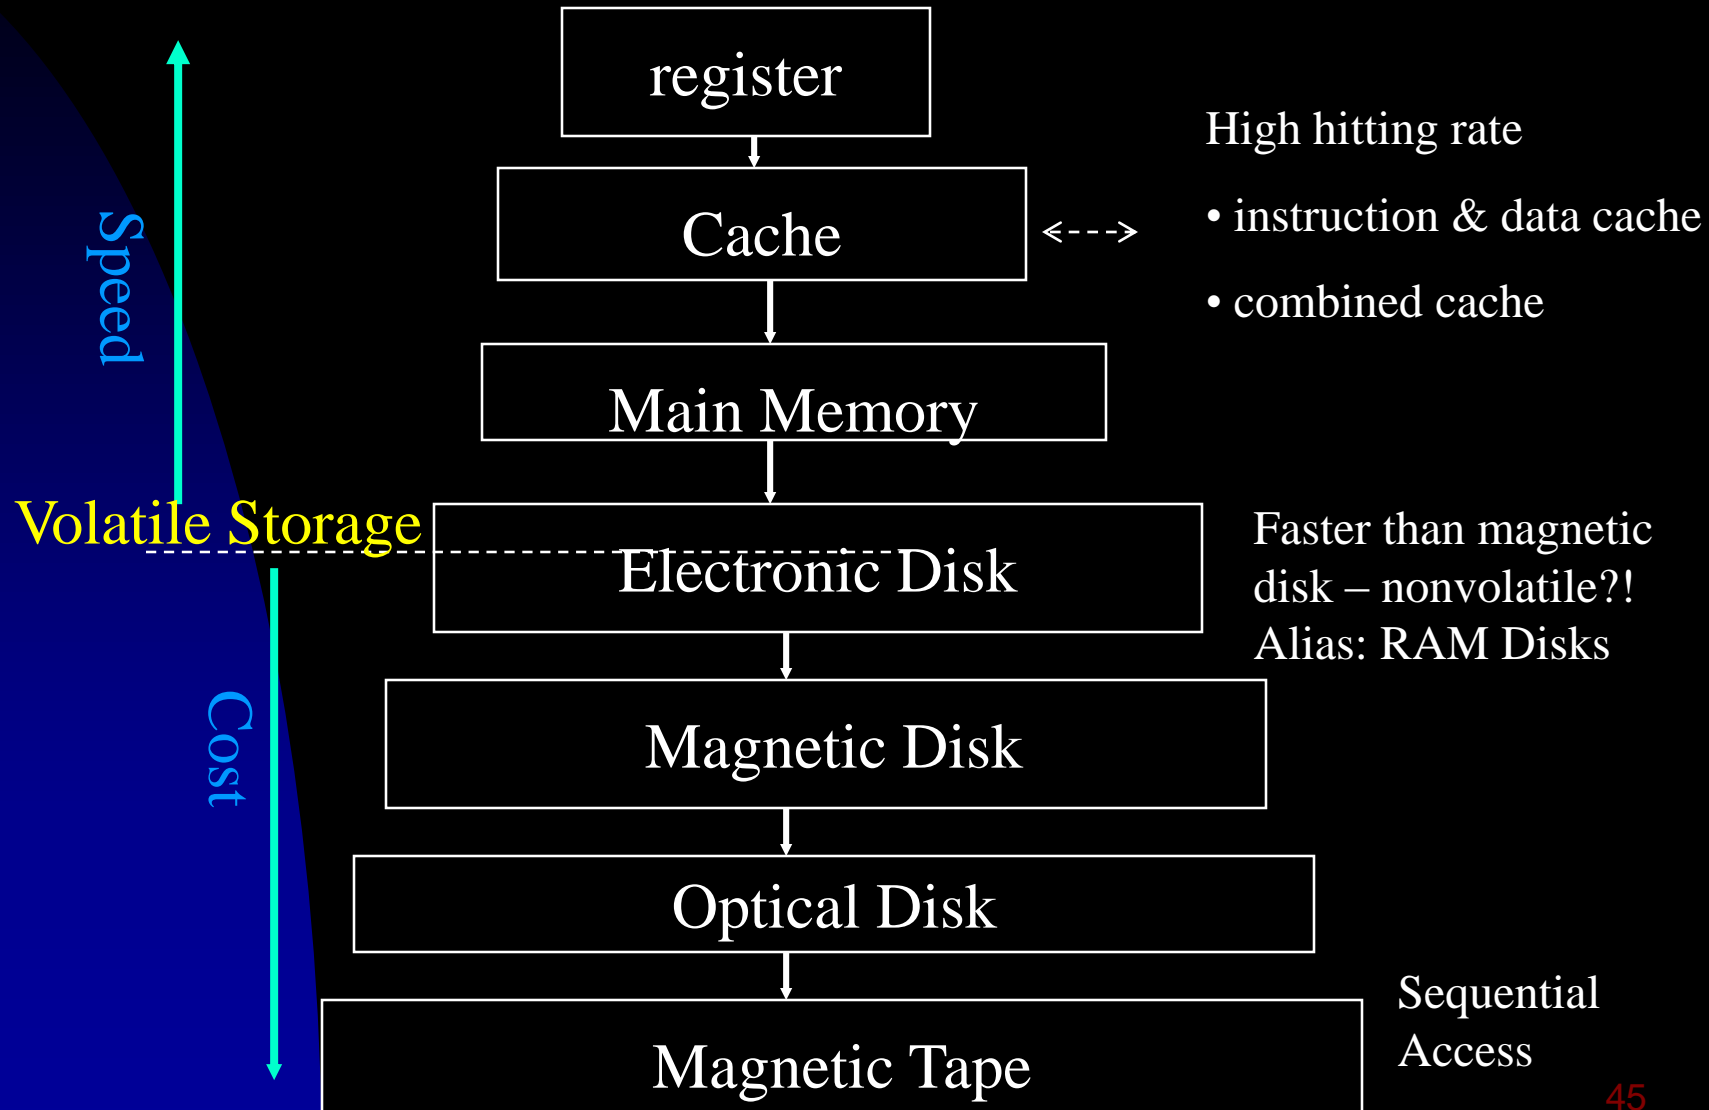
# System Components –Tertiary Storage Devices

- Goals:
  - Backups of disk data, seldom-used data, and long-term archival storage
- Examples:
  - Magnetic tape drives and their tapes, CD & DVD drives and platters.
- Services – OS Supports or Applications' Duty
  - Device mounting and unmounting
  - Exclusive allocation and freeing
  - Data transfers from tertiary devices to secondary storage devices.

# System Components – I/O System Management

- Goal:
  - Hide the peculiarities of specific hardware devices from users
- Components of an I/O System
  - A buffering, caching, and spooling system 解決I/O太慢的問題
  - A general device-driver interface
  - Drivers

# Caching



# Caching





Level	1	2	3	4
Name	Registers	Cache	Memory	Disk
Typical Size	< 1KB	> 16MB	> 16GB	> 100GB
Implementat ion Strategy	Custom memory with multiple ports, CMOS	On-chip or off- chip CMOS SRAM	CMOS DRAM	Magnetic Disks
Access Time (ns)	0.25 – 0.5	0.5 – 2.5	80 – 250	5,000,000
Bandwidth (MB/s)	20,000 – 100,000	5000 – 10,000	1000 – 5000	20 – 150
Managed by	Compiler	Hardware	OS	OS
Backup by	Cache	Memory	Disk	CD/Tape

# Caching

- Caching
  - Information is copied to a faster storage system on a temporary basis
  - Assumption: Data will be used again soon.
    - Programmable registers, instr. cache, etc.
- Cache Management
  - Cache Size and the Replacement Policy
- Movement of Information Between Hierarchy
  - Hardware Design & Controlling Operating Systems

# Caching

- Coherency and Consistency
  - Among several storage levels (vertical)
    - Multitasking vs unitasking
  - Among units of the same storage level , (horizontal), e.g. cache coherency
    - Multiprocessor or distributed systems





# Protection and Security

- Goal
  - Resources are only allowed to be accessed by authorized processes.
- Definitions:
  - Protection – any mechanism for controlling the access of processes or users to the resources defined by the computer system.
  - Security – Defense of a system from external and internal attacks, e.g., viruses, denial of services, etc.

# Protection and Security

- Protected Resources
  - Files, CPU, memory space, etc.
- Protection Services
  - Detection & controlling mechanisms
  - Specification mechanisms
- Distinguishing of Users
  - User names and ID's
  - Group names and GID's
  - Privilege Escalating, e.g., Setuid in Unix
    - To gain extra permissions for an activity.
- Remark: Reliability!

# Kernel Data Structures

- Frequently Used Data Structures
  - Array, List, Stack, Queue, Tree, Hash
  - Bitmaps – A string of  $n$  binary digits to represent the status of  $n$  items.
    - Advantage: Space Efficiency
    - An example is the availability status of disk blocks.

# Computing Environments

- Evolving Environments
  - Transition from the period of scarce resources to the period of ubiquitous access!
  - In the past, portability is achieved by laptops!
    - Remote access is supported in a limited way. Mainframes are prevalent!
  - Now, PC's, mobile devices, and various equipments are connected!
    - High speed networks are available at home and office! Web-computing is popular (e.g., portals).

# Computing Environments

- Mobile Computing
  - Trends: Computing on handheld smartphones and tablets now offers tremendous growth in the wide range of applications, such as email and GPS, augmented-reality applications, but with limitation on screen size, memory/storage capacity, and power/energy consumption.

# Computing Environments

- Distributed/Loosely-Coupled Systems: Heterogeneous or homogeneous computer systems that are networked to provide access to various resources
  - Depend on networking for their functionality
    - Networks vary by the protocols used: TCP/IP, ATM, etc.
  - Be characterized by their node distances
    - Local-area network (LAN)
    - Wide-area network (WAN)
    - Metropolitan-area network (MAN)
    - Personal-area network – distance of few feet

# Computing Environments

- Media – copper wires, fiber strands, satellite wireless transmission, infrared communication, etc.
- Network Operating Systems
  - Autonomous computers
  - A distributed operating system – a single OS controlling the network.

# Computing Environments

- Peer-to-Peer Systems
  - Characteristics: Client and server roles depend on who is requesting or providing a service.
  - Network connectivity is an essential component.
  - Service Availability and Discovery
    - Registration of services: a centralized lookup service or not
    - A discovery protocol
  - Issues:
    - Legal problems in exchanging files.

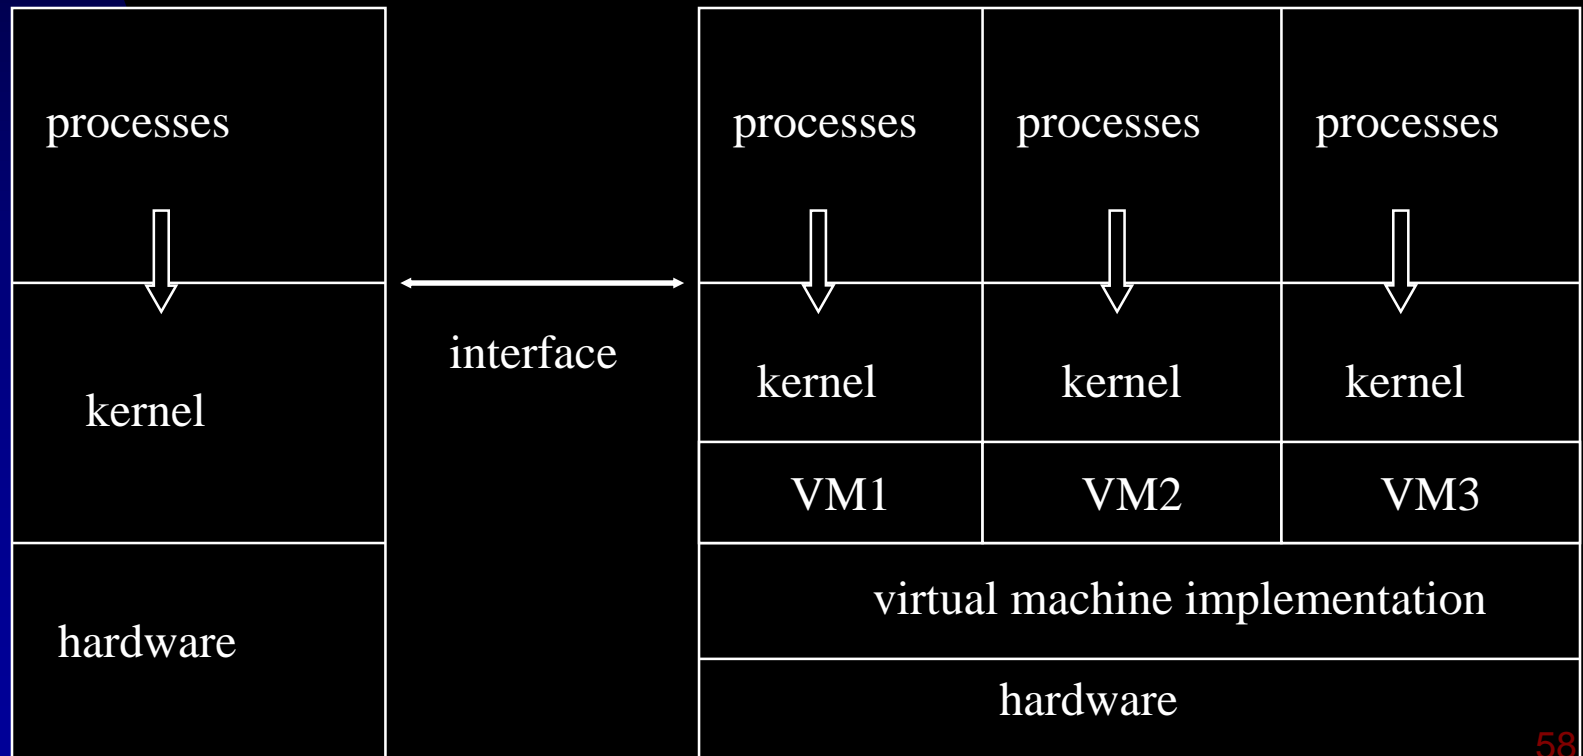


# Computing Environments

- Client-Server Systems
  - Trend: The functionality of clients is improved in the past decades.
  - Categories:
    - Compute-server systems
    - File-server systems

# Computing Environments

- Virtual Machines: provide an interface that is identical to the underlying bare hardware

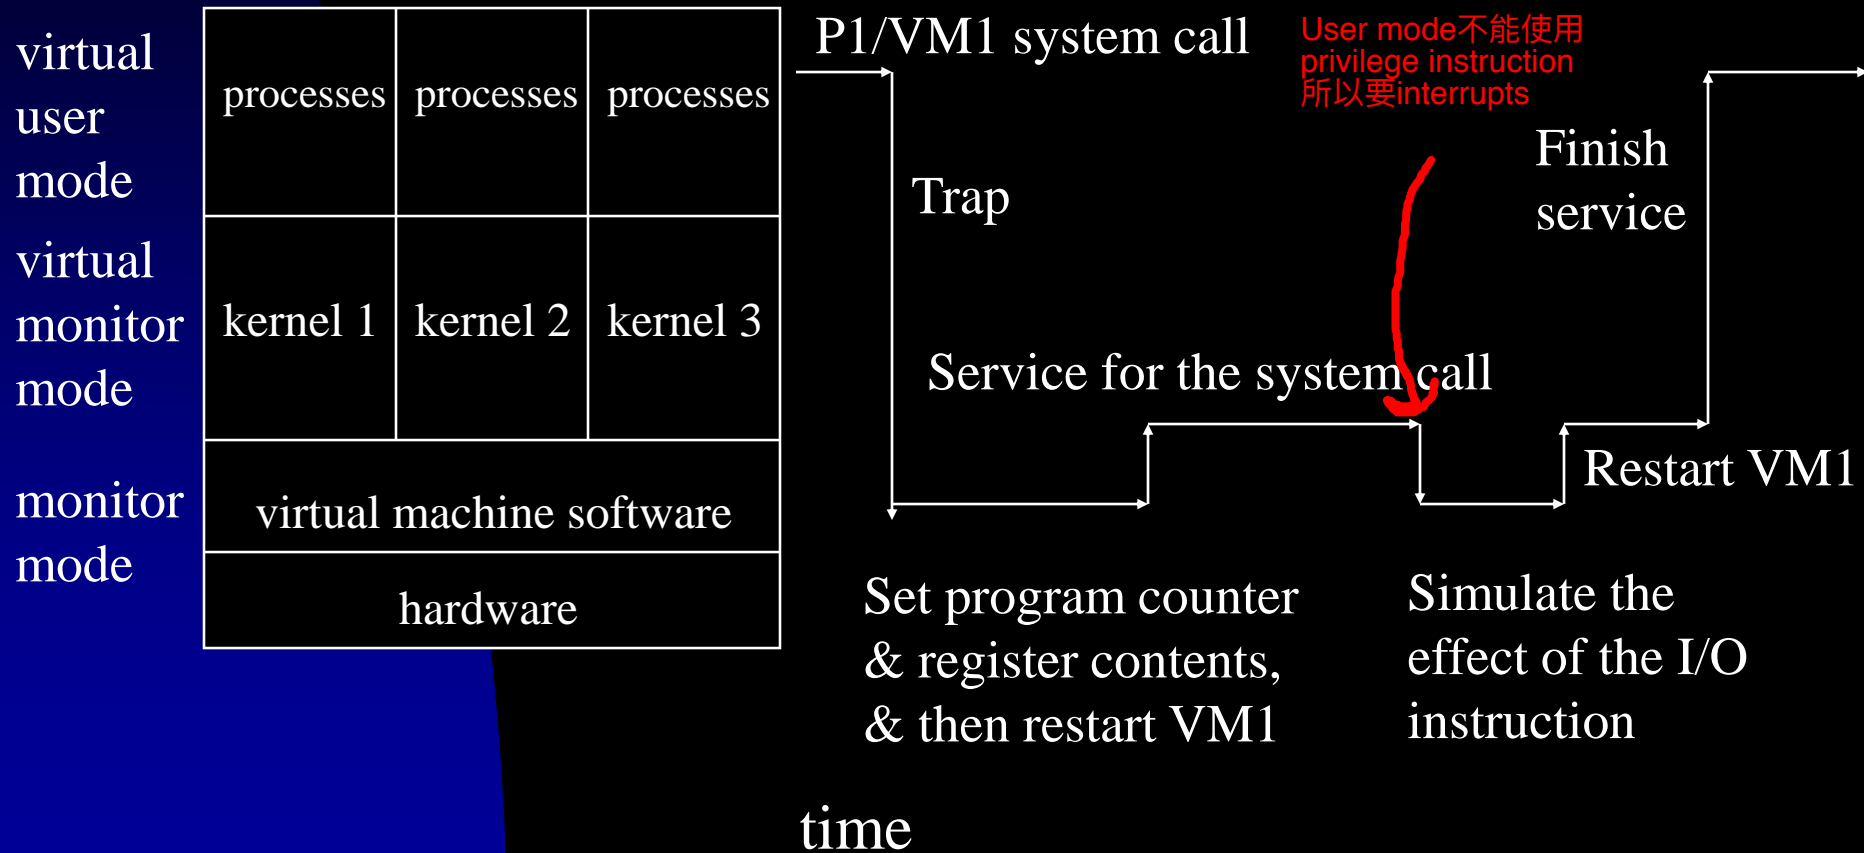


# Computing Environments

- Implementation Issues of Virtual Machines:
  - Emulation of Physical Devices
    - E.g., Disk Systems
      - An IBM minidisk approach
  - User/Monitor Modes
    - (Physical) Monitor Mode
      - Virtual machine software
    - (Physical) User Mode
      - Virtual monitor mode & Virtual user mode

# Computing Environments

## ■ How a Virtual Machine works:



# Computing Environments

- Disadvantages of Virtual Machines:
    - Slow!
      - Execute most instructions directly on the hardware
      - Emulation is slow but is needed for obsolete hardware.
    - No direct sharing of resources
      - Physical devices and communications
- \* I/O could be slow (interpreted) or fast (spooling)

# Computing Environments

- Advantages of Virtual Machines:
  - Complete Protection – Complete Isolation !
  - OS Research & Development
    - System Development Time
  - Extensions to Multiple Personalities, such as Mach (software emulation)
    - Emulations of Machines and OS's, e.g., Windows over Linux
  - System Consolidation

\* Simulation: Programs of a guest system are run on an emulator that translate each of the guest system instructions into the native instruction set of the host system.

# Computing Environments

- Cloud Computing - Delivers computing, storage, and even applications as a service across a network
- Types
  - Public, Private and Hybrid Clouds
  - Software as a service (SaaS), e.g., Gmail.
  - Platform as a service (PaaS), e.g., database server.
  - Infrastructure as a service (IaaS), e.g., storage for backup.

# Computing Environments

- Web-Based Computing
  - Web Technology
    - Portals, network computers, etc.
  - Network connectivity
  - New categories of devices
    - Load balancers
- Embedded Computing
  - Car engines, robots, VCR's, home automation
  - Embedded OS's often have limited features.

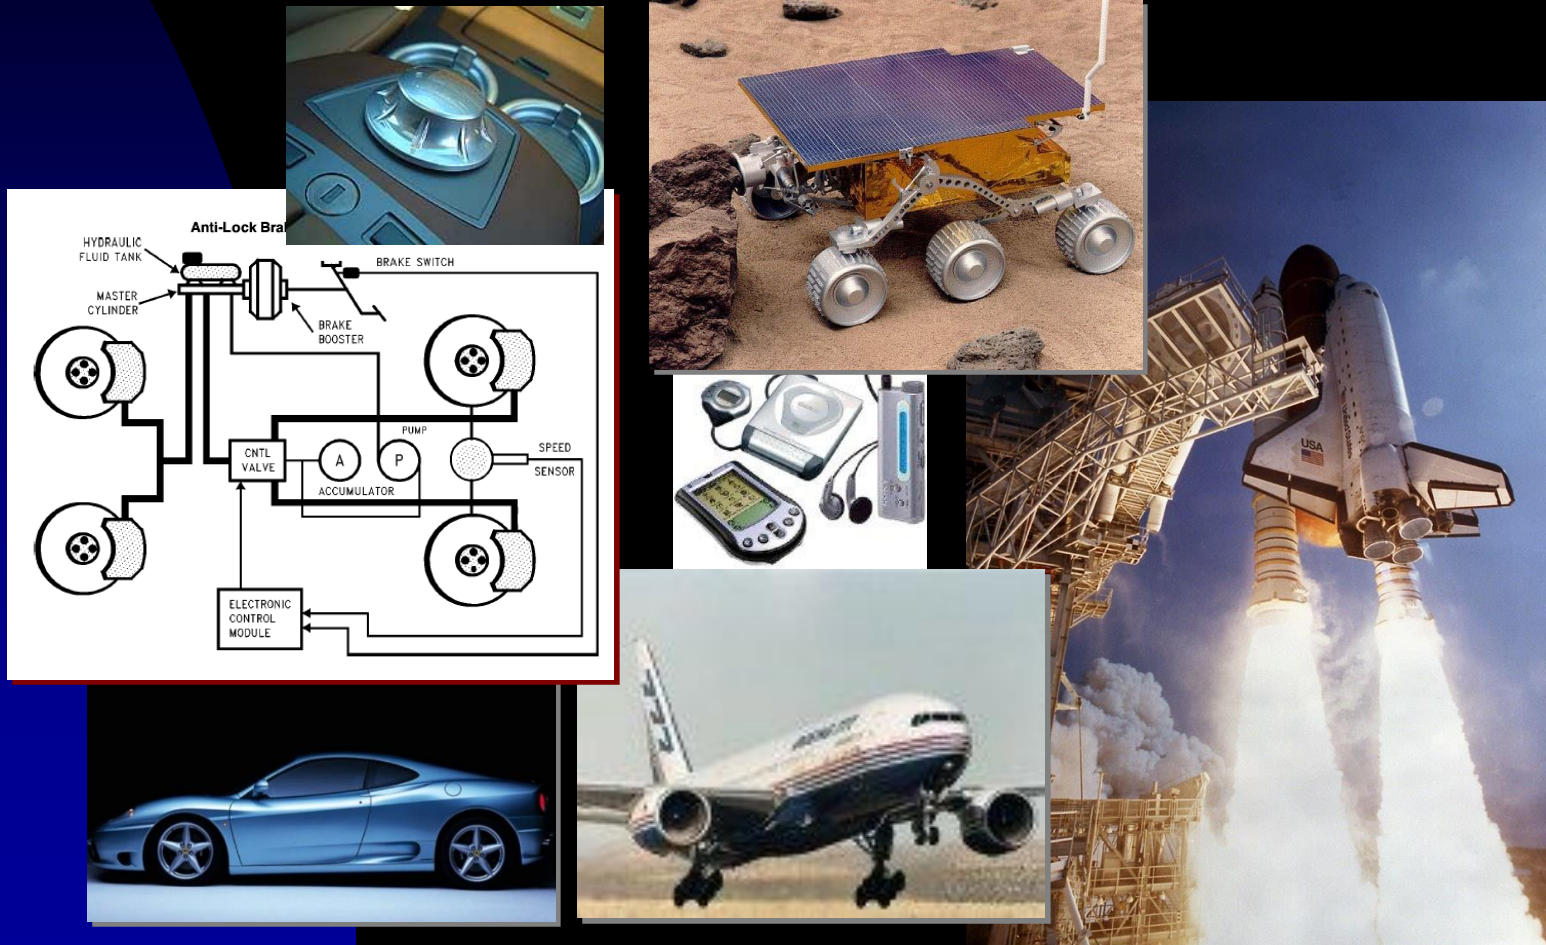


# Computing Environments

- Embedded Computers – Most Prevalent Form of Computers
  - Have a wide variety ranged from car engines to VCR's.
    - General-purpose computers with standard OS's, HW devices with or without embedded OS's
    - Standalone units or members of networks and the Web
  - Tend to have specific tasks and almost always run real-time operating systems.

# Computing Environments

- Real-Time Embedded Computers



# Computing Environments

- Definition: A real-time system is a computer system where a timely response by the computer to external stimuli is vital!
- Hard real-time system: The system has failed if a timing constraint, e.g. deadline, is not met.
  - All delays in the system must be bounded.
  - Many advanced features are absent.

# Computing Environments

- Soft real-time system: Missing a timing constraint is serious but does not necessarily result in a failure unless it is excessive!
  - A critical task has a higher priority.
  - Supported in most commercial OS.
- Real-time means on-time instead of fast

# Open-Source Operating Systems

- Definitions: OS with available source code.
  - Closed-source OS, e.g., MS Windows, or hybrid OS, e.g., iOS.
  - Arguably issues on bugs, security, support, etc.
  - Examples: GNU/Linux, BSD UNIX, and Solaris (up to 2005 versions).